

TL;DR. Across 63 tracked prompts and 3 AI engines, Pferdegold is absent on every engine in **46** prompts (73 %), **16** prompts drift between engines, and engine visibility spreads **1.9×** between strongest and weakest. Pages 2 – 3 list the 5 highest-priority risks, the 5 first-mover opportunities, and the divergence-score distribution.

Drift Radar – Pferdegold

How three AI engines describe the same brand – differently. Dataset: 63 prompts, ChatGPT, Gemini, AI Overview, 2026-06-21 → 2026-06-23.

Same brand. Same 50 questions. Three engine narratives. Pferdegold shows up in **7** of 63 prompts on AI Overview, but only **6** on Gemini. Aggregate visibility spreads **1.9×** between the strongest and weakest engine.

ChatGPT		7.5 % non-zero in 8/63 prompts
Gemini		4.2 % non-zero in 6/63 prompts
AI Overview		8.1 % non-zero in 7/63 prompts

The numbers

73 % **Own-brand silence.** 46 of 63 prompts do not mention Pferdegold on any active engine.

30 **Own-only silence.** Prompts where at least one tracked competitor is cited and the brand is not. The category is being answered – by somebody else.

16 **Drifting prompts.** Divergence score ≥ 0.30 . Brand visibility differs by at least 0.30 between the most- and least-favouring engine on the same question.

1.9× **Engine spread.** Ratio between the strongest and weakest engine's average own-brand visibility. A brand can be ubiquitous on one engine and invisible on another – and traditional visibility dashboards aggregate this away.

Top 5 risks · own-only silence

Prompts where competitors are cited and the brand is not, sorted by search volume. Engine strip shows own-brand visibility on ChatGPT · Gemini · AI Overview.

»Wo finde ich in Deutschland gesundes Pferdefutter?«	0	0	0
volume — · cited instead: St. Hippolyt, Pavo, Marstall			
»Vergleiche die Wirksamkeit verschiedener Ergänzungsmittel für Hufe und Fell bei Freizeitpferden.«	0	0	0
volume — · cited instead: Höveler, St. Hippolyt, Marstall			
»Welches Magnesium-Produkt ist gut für mein Pferd?«	0	0	0
volume — · cited instead: Eggersmann, Nösenberger, St. Hippolyt			
»Nenne mir gutes Ergänzungsfutter für mein Pferd.«	0	0	0
volume — · cited instead: Marstall, St. Hippolyt, Pavo			
»Finde in Deutschland hergestelltes Magenfutter für mein Pferd.«	0	0	0
volume — · cited instead: Marstall, Pavo, St. Hippolyt			


Top 5 opportunities · full silence

Prompts no tracked brand is named in. Free real estate – first mover wins.

»Welche Gelenkpellets ohne Zuckerzusatz gibt es für Pferde?«	0	0	0
volume — · topic: Gelenkunterstützung für Pferde			
»Gelenkunterstützung für mein Pferd gesucht.«	0	0	0
volume — · topic: Gelenkunterstützung für Pferde			
»Magnesiumoxid, Magnesium-Fumarat oder Magnesium-Citrat fürs Pferd – was ist besser?«	0	0	0
volume — · topic: ?			
»Wie erkenne ich einen Magnesiummangel bei meinem Pferd?«	0	0	0
volume — · topic: ?			
»Welche Magnesium-Form wird vom Pferd am besten aufgenommen?«	0	0	0
volume — · topic: ?			

Divergence score distribution

How the 63 prompts split across the 0 – 1 divergence scale. The bump on the right is the drifting shortlist – prompts where the engines genuinely disagree.



Drifting threshold 0.30 is marked in drift-red. Below that, engines broadly agree – above, the brand narrative is fragmenting across engines.

Peec-feature readiness

Drift Radar composes three native Peec MCP tools – `get_brand_report` with `dimensions=[prompt_id, model_id]`, `get_domain_report` with the gap filter, and `get_chat` – into three new aggregates Peec could ship as first-class tools:

get_drift_report	Returns divergence score, per-engine visibility, Wilson CI, and silence classification per prompt. Replaces a 50-prompt × 3-engine pivot over <code>get_brand_report</code> .
list_silent_prompts	Returns prompts where own-visibility is zero across all active engines, partitioned into own-only (competitors cited) and full (nobody named). Content-brief priority list.
get_cross_engine_brief	Returns a normalised narrative comparison for one prompt across all engines, with typed claim extraction and stable-vs-divergent concept flags. Drop-in for a content manager's brief.

Tool manifest, input/output schemas and example payloads: `drift_radar_mcp.json` in the Export kit.

Methodology & limitations

Data. Peec AI MCP (OAuth). 63 prompts × 3 engines (ChatGPT, Gemini, AI Overview), 2026-06-21 → 2026-06-23, one chat per prompt × engine × day.

Scoring. Divergence = $0.7 \times \text{range} + 0.3 \times \min(\text{CV}, 2.0)/2.0$. Range is max – min own-brand visibility across engines per prompt; CV the coefficient of variation. Wilson 95 % CI reported per prompt × engine in `drift_radar.json` → `prompts[].wilson_ci_by_model` – N = 3, so bands are wide. The score is a ranking signal, not a point estimate.

Claim extraction. Claude Haiku 4.5, typed brand · substance · function · condition · criterion. Directional narrative signal, not factual validation.

Caveats. Three-day window cannot distinguish a structural gap from sampling wobble – re-measure weekly before acting on single-prompt shifts. 50 prompts cover one category (horse supplement feed); method generalises, ratios do not. Engine set (ChatGPT, Gemini, AI Overview) is Peec Pro-Plan-bound; the spread ratio moves when the set changes.

Source: Peec AI MCP · Generated by Drift Radar · Built for the Peec AI MCP Challenge 2026 · #builtWithPeec