

**TL;DR.** Across 63 tracked prompts and 3 AI engines, Pferdegold is absent on every engine in **45** prompts (71 %), **17** prompts drift between engines, and engine visibility spreads **1.3×** between strongest and weakest. Pages 2 – 3 list the 5 highest-priority risks, the 5 first-mover opportunities, and the divergence-score distribution.

# Drift Radar – Pferdegold

*How three AI engines describe the same brand – differently. Dataset: 63 prompts, ChatGPT, Gemini, AI Overview, 2026-06-14 → 2026-06-16.*

**Same brand. Same 50 questions. Three engine narratives.** Pferdegold shows up in **7** of 63 prompts on AI Overview, but only **9** on Gemini. Aggregate visibility spreads **1.3×** between the strongest and weakest engine.

|             |                                                                                    |                                          |
|-------------|------------------------------------------------------------------------------------|------------------------------------------|
| ChatGPT     |   | <b>7.9 %</b><br>non-zero in 8/63 prompts |
| Gemini      |   | <b>6.3 %</b><br>non-zero in 9/63 prompts |
| AI Overview |  | <b>8.5 %</b><br>non-zero in 7/63 prompts |

## The numbers

**71 %** **Own-brand silence.** 45 of 63 prompts do not mention Pferdegold on any active engine.

**31** **Own-only silence.** Prompts where at least one tracked competitor is cited and the brand is not. The category is being answered – by somebody else.

**17** **Drifting prompts.** Divergence score  $\geq 0.30$ . Brand visibility differs by at least 0.30 between the most- and least-favouring engine on the same question.

**1.3×** **Engine spread.** Ratio between the strongest and weakest engine's average own-brand visibility. A brand can be ubiquitous on one engine and invisible on another – and traditional visibility dashboards aggregate this away.

## Top 5 risks · own-only silence

Prompts where competitors are cited and the brand is not, sorted by search volume. Engine strip shows own-brand visibility on ChatGPT · Gemini · AI Overview.

|                                                                                                    |   |   |   |
|----------------------------------------------------------------------------------------------------|---|---|---|
| »Wo finde ich in Deutschland gesundes Pferdefutter?«                                               | 0 | 0 | 0 |
| volume — · cited instead: Marstall, St. Hippolyt, Pavo                                             |   |   |   |
| »Gelenkunterstützung für mein Pferd gesucht.«                                                      | 0 | 0 | 0 |
| volume — · cited instead: St. Hippolyt                                                             |   |   |   |
| »Vergleiche die Wirksamkeit verschiedener Ergänzungsmittel für Hufe und Fell bei Freizeitpferden.« | 0 | 0 | 0 |
| volume — · cited instead: Marstall, Pavo                                                           |   |   |   |
| »Welche natürlichen Gelenkpellets für Pferde empfehlst du?«                                        | 0 | 0 | 0 |
| volume — · cited instead: St. Hippolyt                                                             |   |   |   |
| »Welches Magnesium-Produkt ist gut für mein Pferd?«                                                | 0 | 0 | 0 |
| volume — · cited instead: Havens, Marstall, Höveler                                                |   |   |   |


## Top 5 opportunities · full silence

Prompts no tracked brand is named in. Free real estate – first mover wins.

|                                                                                       |   |   |   |
|---------------------------------------------------------------------------------------|---|---|---|
| »Magnesiumoxid, Magnesium-Fumarat oder Magnesium-Citrat fürs Pferd – was ist besser?« | 0 | 0 | 0 |
| volume — · topic: ?                                                                   |   |   |   |
| »Wie erkenne ich einen Magnesiummangel bei meinem Pferd?«                             | 0 | 0 | 0 |
| volume — · topic: ?                                                                   |   |   |   |
| »Welche Magnesium-Form wird vom Pferd am besten aufgenommen?«                         | 0 | 0 | 0 |
| volume — · topic: ?                                                                   |   |   |   |
| »Finde natürliche Pellets zur Unterstützung der Pferdegesundheit.«                    | 0 | 0 | 0 |
| volume — · topic: Pferde Nahrungsergänzung                                            |   |   |   |
| »Brauche ich für mein nervöses Sportpferd Magnesium pur oder ein Tryptophan-Komplex?« | 0 | 0 | 0 |
| volume — · topic: ?                                                                   |   |   |   |

## Divergence score distribution

How the 63 prompts split across the 0 – 1 divergence scale. The bump on the right is the drifting shortlist – prompts where the engines genuinely disagree.



Drifting threshold 0.30 is marked in drift-red. Below that, engines broadly agree – above, the brand narrative is fragmenting across engines.

## Peec-feature readiness

Drift Radar composes three native Peec MCP tools – `get_brand_report` with `dimensions=[prompt_id, model_id]`, `get_domain_report` with the gap filter, and `get_chat` – into three new aggregates Peec could ship as first-class tools:

|                        |                                                                                                                                                                                        |
|------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| get_drift_report       | Returns divergence score, per-engine visibility, Wilson CI, and silence classification per prompt. Replaces a 50-prompt × 3-engine pivot over get_brand_report.                        |
| list_silent_prompts    | Returns prompts where own-visibility is zero across all active engines, partitioned into own-only (competitors cited) and full (nobody named). Content-brief priority list.            |
| get_cross_engine_brief | Returns a normalised narrative comparison for one prompt across all engines, with typed claim extraction and stable-vs-divergent concept flags. Drop-in for a content manager's brief. |

Tool manifest, input/output schemas and example payloads: `drift_radar_mcp.json` in the Export kit.

## Methodology & limitations

**Data.** Peec AI MCP (OAuth). 63 prompts × 3 engines (ChatGPT, Gemini, AI Overview), 2026-06-14 → 2026-06-16, one chat per prompt × engine × day.

**Scoring.** Divergence =  $0.7 \times \text{range} + 0.3 \times \min(\text{CV}, 2.0)/2.0$ . Range is max – min own-brand visibility across engines per prompt; CV the coefficient of variation. Wilson 95 % CI reported per prompt × engine in `drift_radar.json` → `prompts[].wilson_ci_by_model` – N = 3, so bands are wide. The score is a ranking signal, not a point estimate.

**Claim extraction.** Claude Haiku 4.5, typed brand · substance · function · condition · criterion. Directional narrative signal, not factual validation.

**Caveats.** Three-day window cannot distinguish a structural gap from sampling wobble – re-measure weekly before acting on single-prompt shifts. 50 prompts cover one category (horse supplement feed); method generalises, ratios do not. Engine set (ChatGPT, Gemini, AI Overview) is Peec Pro-Plan-bound; the spread ratio moves when the set changes.

Source: Peec AI MCP · Generated by Drift Radar · Built for the Peec AI MCP Challenge 2026 · #builtWithPeec